

A Hacker's Guide to NCURSES

Contents

- [Abstract](#)
- [Objective of the Package](#)
 - [Why System V Curses?](#)
 - [How to Design Extensions](#)
- [Portability and Configuration](#)
- [Documentation Conventions](#)
- [How to Report Bugs](#)
- [A Tour of the Ncurses Library](#)
 - [Library Overview](#)
 - [The Engine Room](#)
 - [Keyboard Input](#)
 - [Mouse Events](#)
 - [Output and Screen Updating](#)
- [The Forms and Menu Libraries](#)
- [A Tour of the Terminfo Compiler](#)
 - [Translation of Non-**use** Capabilities](#)
 - [Use Capability Resolution](#)
 - [Source-Form Translation](#)
- [Other Utilities](#)
- [Style Tips for Developers](#)
- [Porting Hints](#)

Abstract

This document is a hacker's tour of the **ncurses** library and utilities. It discusses design philosophy, implementation methods, and the conventions used for coding and documentation. It is recommended reading for anyone who is interested in porting, extending or improving the package.

Objective of the Package

The objective of the **ncurses** package is to provide a free software API for character-cell terminals and terminal emulators with the following characteristics:

- Source-compatible with historical curses implementations (including the original BSD curses and System V curses).
- Conformant with the XSI Curses standard issued as part of XPG4 by X/Open.
- High-quality — stable and reliable code, wide portability, good packaging, superior documentation.
- Featureful — should eliminate as much of the drudgery of C interface programming as possible, freeing programmers to think at a higher level of design.

These objectives are in priority order. So, for example, source compatibility with older version must trump featurefulness — we cannot add features if it means breaking the portion of the API corresponding to historical curses versions.

Why System V Curses?

We used System V curses as a model, reverse-engineering their API, in order to fulfill the first two objectives.

System V curses implementations can support BSD curses programs with just a recompilation, so by capturing the System V API we also capture BSD's.

More importantly for the future, the XSI Curses standard issued by X/Open is explicitly and closely modeled on System V. So conformance with System V took us most of the way to base-level XSI conformance.

How to Design Extensions

The third objective (standards conformance) requires that it be easy to condition source code using **ncurses** so that the absence of nonstandard extensions does not break the code.

Accordingly, we have a policy of associating with each nonstandard extension a feature macro, so that ncurses client code can use this macro to condition in or out the code that requires the **ncurses** extension.

For example, there is a macro `NCURSES_MOUSE_VERSION` which XSI Curses does not define, but which is defined in the **ncurses** library header. You can use this to condition the calls to the mouse API calls.

Portability and Configuration

Code written for **ncurses** may assume an ANSI-standard C compiler and POSIX-compatible OS interface. It may also assume the presence of a System-V-compatible *select(2)* call.

We encourage (but do not require) developers to make the code friendly to less-capable UNIX environments wherever possible.

We encourage developers to support OS-specific optimizations and methods not available under POSIX/ANSI, provided only that:

- All such code is properly conditioned so the build process does not attempt to compile it under a plain ANSI/POSIX environment.
- Adding such implementation methods does not introduce incompatibilities in the **ncurses** API between platforms.

We use GNU `autoconf(1)` as a tool to deal with portability issues. The right way to leverage an OS-specific feature is to modify the autoconf specification files (`configure.in` and `aclocal.m4`) to set up a new feature macro, which you then use to condition your code.

Documentation Conventions

There are three kinds of documentation associated with this package. Each has a different preferred format:

- Package-internal files (README, INSTALL, TO-DO etc.)
- Manual pages.
- Everything else (i.e., narrative documentation).

Our conventions are simple:

1. **Maintain package-internal files in plain text.** The expected viewer for them *more(1)* or an editor window; there is no point in elaborate mark-up.
2. **Mark up manual pages in the man macros.** These have to be viewable through traditional *man(1)* programs.
3. **Write everything else in HTML.**

When in doubt, HTMLize a master and use *lynx(1)* to generate plain ASCII (as we do for the announcement document).

The reason for choosing HTML is that it is (a) well-adapted for on-line browsing through viewers that are everywhere; (b) more easily readable as plain text than most other mark-ups, if you do not have a viewer; and (c) carries enough information that you can generate a nice-looking printed version from it. Also, of course, it make exporting things like the announcement document to WWW pretty trivial.

How to Report Bugs

The reporting address for bugs is bug-ncurses@gnu.org. This is a majordomo list; to join, write to bug-ncurses-request@gnu.org with a message containing the line:

```
subscribe <name>@<host.domain>
```

The `ncurses` code is maintained by a small group of volunteers. While we try our best to fix bugs promptly, we simply do not have a lot of hours to spend on elementary hand-holding. We rely on intelligent cooperation from our users. If you think you have found a bug in `ncurses`, there are some steps you can take before contacting us that will help get the bug fixed quickly.

In order to use our bug-fixing time efficiently, we put people who show us they have taken these steps at the head of our queue. This means that if you do not, you will probably end up at the tail end and have to wait a while.

1. Develop a recipe to reproduce the bug.

Bugs we can reproduce are likely to be fixed very quickly, often within days. The most effective single thing you can do to get a quick fix is develop a way we can duplicate the bad behavior — ideally, by giving us source for a small, portable test program that breaks the library. (Even better is a keystroke recipe using one of the test programs provided with the distribution.)

2. Try to reproduce the bug on a different terminal type.

In our experience, most of the behaviors people report as library bugs are actually due to subtle problems in terminal descriptions. This is especially likely to be true if you are using a traditional asynchronous terminal or PC-based terminal emulator, rather than xterm or a UNIX console entry.

It is therefore extremely helpful if you can tell us whether or not your problem reproduces on other terminal types. Usually you will have both a console type and xterm available; please tell us whether or not your bug reproduces on both.

If you have xterm available, it is also good to collect xterm reports for different window sizes. This is especially true if you normally use an unusual xterm window size — a surprising number of the bugs we have seen are either triggered or masked by these.

3. Generate and examine a trace file for the broken behavior.

Recompile your program with the debugging versions of the libraries. Insert a `trace()` call with the argument set to `TRACE_UPDATE`. (See ["Writing Programs with NCURSES"](#) for details on trace levels.) Reproduce your bug, then look at the trace file to see what the library was actually doing.

Another frequent cause of apparent bugs is application coding errors that cause the wrong things to be put on the virtual screen. Looking at the virtual-screen dumps in the trace file will tell you immediately if this is happening, and save you from the possible embarrassment of being told that the bug is in your code and is your problem rather than ours.

If the virtual-screen dumps look correct but the bug persists, it is possible to crank up the trace level to give more and more information about the library's update actions and the control sequences it issues to perform them. The test directory of the distribution contains a tool for digesting these logs to make them less tedious to wade through.

Often you will find terminfo problems at this stage by noticing that the escape sequences put out for various capabilities are wrong. If not, you are likely to learn enough to be able to characterize any bug in the screen-update logic quite exactly.

4. Report details and symptoms, not just interpretations.

If you do the preceding two steps, it is very likely that you will discover the nature of the problem yourself and be able to send us a fix. This will create happy feelings all around and earn you good karma for the first time you run into a bug you really cannot characterize and fix yourself.

If you are still stuck, at least you will know what to tell us. Remember, we need details. If you guess about what is safe to leave out, you are too likely to be wrong.

If your bug produces a bad update, include a trace file. Try to make the trace at the *least* voluminous level that pins down the bug. Logs that have been through tracemunch are OK, it does not throw away any information (actually

they are better than un-munched ones because they are easier to read).

If your bug produces a core-dump, please include a symbolic stack trace generated by `gdb(1)` or your local equivalent.

Tell us about every terminal on which you have reproduced the bug — and every terminal on which you cannot. Ideally, send us terminfo sources for all of these (yours might differ from ours).

Include your ncurses version and your OS/machine type, of course! You can find your ncurses version in the `curses.h` file.

If your problem smells like a logic error or in cursor movement or scrolling or a bad capability, there are a couple of tiny test frames for the library algorithms in the `progs` directory that may help you isolate it. These are not part of the normal build, but do have their own make productions.

The most important of these is `mvcur`, a test frame for the cursor-movement optimization code. With this program, you can see directly what control sequences will be emitted for any given cursor movement or scroll/insert/delete operations. If you think you have got a bad capability identified, you can disable it and test again. The program is command-driven and has on-line help.

If you think the vertical-scroll optimization is broken, or just want to understand how it works better, build `hashmap` and read the header comments of `hardscroll.c` and `hashmap.c`; then try it out. You can also test the hardware-scrolling optimization separately with `hardscroll`.

A Tour of the Ncurses Library

Library Overview

Most of the library is superstructure — fairly trivial convenience interfaces to a small set of basic functions and data structures used to manipulate the virtual screen (in particular, none of this code does any I/O except through calls to more fundamental modules described below). The files

```
lib_addch.c lib_bkgd.c lib_box.c lib_chgat.c lib_clear.c lib_clearok.c lib_clrbot.c
lib_clreol.c lib_colorset.c lib_data.c lib_delch.c lib_delwin.c lib_echo.c
lib_erase.c lib_gen.c lib_getstr.c lib_hline.c lib_immedok.c lib_inchstr.c
lib_insch.c lib_insdel.c lib_insstr.c lib_instr.c lib_isendwin.c lib_keyname.c
lib_leaveok.c lib_move.c lib_mvwin.c lib_overlay.c lib_pad.c lib_printw.c
lib_redrawln.c lib_scanw.c lib_screen.c lib_scroll.c lib_scrollok.c lib_scrreg.c
lib_set_term.c lib_slk.c lib_slkatr_set.c lib_slkatrof.c lib_slkatron.c
lib_slkatrset.c lib_slkattrib.c lib_slkclear.c lib_slkcolor.c lib_slkinit.c
lib_slklab.c lib_slkrefr.c lib_slkset.c lib_slktouch.c lib_touch.c lib_unctrl.c
lib_vline.c lib_wattroff.c lib_wattron.c lib_window.c
```

are all in this category. They are very unlikely to need change, barring bugs or some fundamental reorganization in the underlying data structures.

These files are used only for debugging support:

```
lib_trace.c lib_traceatr.c lib_tracebits.c lib_tracechr.c lib_tracedmp.c  
lib_tracemse.c trace_buf.c
```

It is rather unlikely you will ever need to change these, unless you want to introduce a new debug trace level for some reason.

There is another group of files that do direct I/O via *tputs()*, computations on the terminal capabilities, or queries to the OS environment, but nevertheless have only fairly low complexity. These include:

```
lib_acs.c lib_beep.c lib_color.c lib_endwin.c lib_initscr.c lib_longname.c  
lib_newterm.c lib_options.c lib_termcap.c lib_ti.c lib_tparm.c lib_tputs.c  
lib_vidattr.c read_entry.c.
```

They are likely to need revision only if ncurses is being ported to an environment without an underlying terminfo capability representation.

These files have serious hooks into the tty driver and signal facilities:

```
lib_kernel.c lib_baudrate.c lib_raw.c lib_tstp.c lib_twait.c
```

If you run into porting snafus moving the package to another UNIX, the problem is likely to be in one of these files. The file `lib_print.c` uses `sleep(2)` and also falls in this category.

Almost all of the real work is done in the files

```
hardscroll.c hashmap.c lib_addch.c lib_doupdate.c lib_getch.c lib_mouse.c  
lib_mvcur.c lib_refresh.c lib_setup.c lib_vidattr.c
```

Most of the algorithmic complexity in the library lives in these files. If there is a real bug in **ncurses** itself, it is probably here. We will tour some of these files in detail below (see [The Engine Room](#)).

Finally, there is a group of files that is actually most of the terminfo compiler. The reason this code lives in the **ncurses** library is to support fallback to `/etc/termcap`. These files include

```
alloc_entry.c captainfo.c comp_captab.c comp_error.c comp_hash.c comp_parse.c  
comp_scan.c parse_entry.c read_termcap.c write_entry.c
```

We will discuss these in the compiler tour.

The Engine Room

Keyboard Input

All ncurses input funnels through the function `wgetch()`, defined in `lib_getch.c`. This function is tricky; it has to poll for keyboard and mouse events and do a running match of incoming input against the set of defined special keys.

The central data structure in this module is a FIFO queue, used to match multiple-character input sequences against special-key capabilities; also to implement pushback via `ungetch()`.

The `wgetch()` code distinguishes between function key sequences and the same sequences typed manually by doing a timed wait after each input character that could lead a function key sequence. If the entire sequence takes less than 1 second, it is assumed to have been generated by a function key press.

Hackers bruised by previous encounters with variant `select(2)` calls may find the code in `lib_twait.c` interesting. It deals with the problem that some BSD selects do not return a reliable time-left value. The function `timed_wait()` effectively simulates a System V `select`.

Mouse Events

If the mouse interface is active, `wgetch()` polls for mouse events each call, before it goes to the keyboard for input. It is up to `lib_mouse.c` how the polling is accomplished; it may vary for different devices.

Under `xterm`, however, mouse event notifications come in via the keyboard input stream. They are recognized by having the **`kmous`** capability as a prefix. This is kind of klugey, but trying to wire in recognition of a mouse key prefix without going through the function-key machinery would be just too painful, and this turns out to imply having the prefix somewhere in the function-key capabilities at terminal-type initialization.

This kluge only works because **`kmous`** is not actually used by any historic terminal type or curses implementation we know of. Best guess is it is a relic of some forgotten experiment in-house at Bell Labs that did not leave any traces in the publicly-distributed System V terminfo files. If System V or XPG4 ever gets serious about using it again, this kluge may have to change.

Here are some more details about mouse event handling:

The `lib_mouse()` code is logically split into a lower level that accepts event reports in a device-dependent format and an upper level that parses mouse gestures and filters events. The mediating data structure is a circular queue of event structures.

Functionally, the lower level's job is to pick up primitive events and put them on the circular queue. This can happen in one of two ways: either (a) `_nc_mouse_event()` detects a series of incoming mouse reports and queues them, or (b) code in `lib_getch.c` detects the **`kmous`** prefix in the keyboard input stream and calls `_nc_mouse_inline` to queue up a series of adjacent mouse reports.

In either case, `_nc_mouse_parse()` should be called after the series is accepted to parse the digested mouse reports (low-level events) into a gesture (a high-level or composite event).

Output and Screen Updating

With the single exception of character echoes during a `wgetnstr()` call (which simulates cooked-mode line editing in an ncurses window), the library normally does all its output at refresh time.

The main job is to go from the current state of the screen (as represented in the `curscr` window structure) to the desired new state (as represented in the `newscr` window structure), while doing as little I/O as possible.

The brains of this operation are the modules `hashmap.c`, `hardscroll.c` and `lib_douupdate.c`; the latter two use `lib_mvcur.c`. Essentially, what happens looks like this:

- The `hashmap.c` module tries to detect vertical motion changes between the real and virtual screens. This information is represented by the `oldindex` members in the `newscr` structure. These are modified by `vertical-motion` and `clear` operations, and both are re-initialized after each update. To this change-journaling information, the `hashmap` code adds deductions made using a modified Heckel algorithm on hash values generated from the line contents.
- The `hardscroll.c` module computes an optimum set of scroll, insertion, and deletion operations to make the indices match. It calls `_nc_mvcur_scrolln()` in `lib_mvcur.c` to do those motions.
- Then `lib_douupdate.c` goes to work. Its job is to do line-by-line transformations of `curscr` lines to `newscr` lines. Its main tool is the routine `mvcur()` in `lib_mvcur.c`. This routine does cursor-movement optimization, attempting to get from given screen location A to given location B in the fewest output characters possible.

If you want to work on screen optimizations, you should use the fact that (in the trace-enabled version of the library) enabling the `TRACE_TIMES` trace level causes a report to be emitted after each screen update giving the elapsed time and a count of characters emitted during the update. You can use this to tell when an update optimization improves efficiency.

In the trace-enabled version of the library, it is also possible to disable and re-enable various optimizations at runtime by tweaking the variable `_nc_optimize_enable`. See the file `include/curses.h.in` for mask values, near the end.

The Forms and Menu Libraries

The forms and menu libraries should work reliably in any environment you can port `ncurses` to. The only portability issue anywhere in them is what flavor of regular expressions the built-in form field type `TYPE_REGEX` will recognize.

The configuration code prefers the POSIX regex facility, modeled on System V's, but will settle for BSD regexps if the former is not available.

Historical note: the panels code was written primarily to assist in porting `u386mon 2.0` (comp.sources.misc v14i001-4) to systems lacking panels support; `u386mon 2.10` and beyond use it. This version has been slightly cleaned up for `ncurses`.

A Tour of the Terminfo Compiler

The **`ncurses`** implementation of **`tic`** is rather complex internally; it has to do a trying combination of missions. This starts with the fact that, in addition to its normal duty of compiling terminfo sources into loadable terminfo binaries, it has to be able to handle termcap syntax and compile that too into terminfo entries.

The implementation therefore starts with a table-driven, dual-mode lexical analyzer (in `comp_scan.c`). The lexer chooses its mode (termcap or terminfo) based on the first `“,”` or `“:”` it finds in each entry. The lexer does all the work of recognizing capability

names and values; the grammar above it is trivial, just "parse entries till you run out of file".

Translation of Non-use Capabilities

Translation of most things besides **use** capabilities is pretty straightforward. The lexical analyzer's tokenizer hands each capability name to a hash function, which drives a table lookup. The table entry yields an index which is used to look up the token type in another table, and controls interpretation of the value.

One possibly interesting aspect of the implementation is the way the compiler tables are initialized. All the tables are generated by various awk/sed/sh scripts from a master table `include/Caps`; these scripts actually write C initializers which are linked to the compiler. Furthermore, the hash table is generated in the same way, so it doesn't have to be generated at compiler startup time (another benefit of this organization is that the hash table can be in shareable text space).

Thus, adding a new capability is usually pretty trivial, just a matter of adding one line to the `include/Caps` file. We will have more to say about this in the section on [Source-Form Translation](#).

Use Capability Resolution

The background problem that makes **tic** tricky is not the capability translation itself, it is the resolution of **use** capabilities. Older versions would not handle forward **use** references for this reason (that is, a using terminal always had to follow its use target in the source file). By doing this, they got away with a simple implementation tactic; compile everything as it blows by, then resolve uses from compiled entries.

This will not do for **ncurses**. The problem is that that the whole compilation process has to be embeddable in the **ncurses** library so that it can be called by the startup code to translate termcap entries on the fly. The embedded version cannot go promiscuously writing everything it translates out to disk — for one thing, it will typically be running with non-root permissions.

So our **tic** is designed to parse an entire terminfo file into a doubly-linked circular list of entry structures in-core, and then do **use** resolution in-memory before writing everything out. This design has other advantages: it makes forward and back use-references equally easy (so we get the latter for free), and it makes checking for name collisions before they are written out easy to do.

And this is exactly how the embedded version works. But the stand-alone user-accessible version of **tic** partly reverts to the historical strategy; it writes to disk (not keeping in core) any entry with no **use** references.

This is strictly a core-economy kluge, implemented because the terminfo master file is large enough that some core-poor systems swap like crazy when you compile it all in memory...there have been reports of this process taking **three hours**, rather than the twenty seconds or less typical on the author's development box.

So. The executable **tic** passes the entry-parser a hook that *immediately* writes out the referenced entry if it has no use capabilities. The compiler main loop refrains

from adding the entry to the in-core list when this hook fires. If some other entry later needs to reference an entry that got written immediately, that is OK; the resolution code will fetch it off disk when it cannot find it in core.

Name collisions will still be detected, just not as cleanly. The `write_entry()` code complains before overwriting an entry that postdates the time of **tic**'s first call to `write_entry()`. Thus it will complain about overwriting entries newly made during the **tic** run, but not about overwriting ones that predate it.

Source-Form Translation

Another use of **tic** is to do source translation between various termcap and terminfo formats. There are more variants out there than you might think; the ones we know about are described in the **captoinfo(1)** manual page.

The translation output code (`dump_entry()` in `ncurses/dump_entry.c`) is shared with the **infocmp(1)** utility. It takes the same internal representation used to generate the binary form and dumps it to standard output in a specified format.

The `include/Caps` file has a header comment describing ways you can specify source translations for nonstandard capabilities just by altering the master table. It is possible to set up capability aliasing or tell the compiler to plain ignore a given capability without writing any C code at all.

For circumstances where you need to do algorithmic translation, there are functions in `parse_entry.c` called after the parse of each entry that are specifically intended to encapsulate such translations. This, for example, is where the AIX **box1** capability get translated to an **acsc** string.

Other Utilities

The **infocmp** utility is just a wrapper around the same entry-dumping code used by **tic** for source translation. Perhaps the one interesting aspect of the code is the use of a predicate function passed in to `dump_entry()` to control which capabilities are dumped. This is necessary in order to handle both the ordinary De-compilation case and entry difference reporting.

The **tput** and **clear** utilities just do an entry load followed by a `tputs()` of a selected capability.

Style Tips for Developers

See the TO-DO file in the top-level directory of the source distribution for additions that would be particularly useful.

The prefix `_nc_` should be used on library public functions that are not part of the curses API in order to prevent pollution of the application namespace. If you have to add to or modify the function prototypes in `curses.h.in`, read `ncurses/MKlib_gen.sh` first so you can avoid breaking XSI conformance. Please join the ncurses mailing list. See the INSTALL file in the top level of the distribution for details on the list.

Look for the string `FIXME` in source files to tag minor bugs and potential problems that could use fixing.

Do not try to auto-detect OS features in the main body of the C code. That is the job of the configuration system.

To hold down complexity, do make your code data-driven. Especially, if you can drive logic from a table filtered out of `include/Caps`, do it. If you find you need to augment the data in that file in order to generate the proper table, that is still preferable to ad-hoc code — that is why the fifth field (`flags`) is there.

Have fun!

Porting Hints

The following notes are intended to be a first step towards DOS and Macintosh ports of the ncurses libraries.

The following library modules are “pure curses”; they operate only on the curses internal structures, do all output through other curses calls (not including `tputs()` and `putp()`) and do not call any other UNIX routines such as `signal(2)` or the `stdio` library. Thus, they should not need to be modified for single-terminal ports.

```
lib_addch.c lib_addstr.c lib_bkgd.c lib_box.c lib_clear.c lib_clrbot.c lib_clreol.c
lib_delch.c lib_delwin.c lib_erase.c lib_inchstr.c lib_insch.c lib_insdcl.c
lib_insstr.c lib_keyname.c lib_move.c lib_mvwin.c lib_newwin.c lib_overlay.c
lib_pad.c lib_printw.c lib_refresh.c lib_scanw.c lib_scroll.c lib_scrreg.c
lib_set_term.c lib_touch.c lib_tparm.c lib_tputs.c lib_unctrl.c lib_window.c
panel.c
```

This module is pure curses, but calls `outstr()`:

```
lib_getstr.c
```

These modules are pure curses, except that they use `tputs()` and `putp()`:

```
lib_beep.c lib_color.c lib_endwin.c lib_options.c lib_slk.c lib_vidattr.c
```

This modules assist in POSIX emulation on non-POSIX systems:

```
sigaction.c
signal calls
```

The following source files will not be needed for a single-terminal-type port.

```
alloc_entry.c captainfo.c clear.c comp_captab.c comp_error.c comp_hash.c
comp_main.c comp_parse.c comp_scan.c dump_entry.c infocmp.c parse_entry.c
read_entry.c tput.c write_entry.c
```

The following modules will use `open()/read()/write()/close()/lseek()` on files, but no other OS calls.

```
lib_screen.c
used to read/write screen dumps
```

`lib_trace.c`
used to write trace data to the logfile

Modules that would have to be modified for a port start here:

The following modules are “pure curses” but contain assumptions inappropriate for a memory-mapped port.

`lib_longname.c`
assumes there may be multiple terminals
`lib_acs.c`
assumes `acs_map` as a double indirection
`lib_mvcur.c`
assumes cursor moves have variable cost
`lib_termcap.c`
assumes there may be multiple terminals
`lib_ti.c`
assumes there may be multiple terminals

The following modules use UNIX-specific calls:

`lib_douupdate.c`
input checking
`lib_getch.c`
`read()`
`lib_initscr.c`
`getenv()`
`lib_newterm.c`
`lib_baudrate.c`
`lib_kernel.c`
various tty-manipulation and system calls
`lib_raw.c`
various tty-manipulation calls
`lib_setup.c`
various tty-manipulation calls
`lib_restart.c`
various tty-manipulation calls
`lib_tstp.c`
signal-manipulation calls
`lib_twait.c`
`gettimeofday()`, `select()`.

Eric S. Raymond <esr@snark.thyrsus.com>

(Note: This is *not* the [bug address](#)!)